

Computer Organization & Architecture

Performance evaluation

- For a non-pipelined processor, execution time, T , of a program that has a dynamic instruction count of I_c is:

$$T = I_c \times S / f$$

where S is average number of clock cycles it takes to fetch and execute one instruction, and f is clock rate in cycles per second.

- For non-pipelined execution, throughput, P_{np} , is:

$$P_{np} = f / S$$

- For a processor that uses five cycles to execute all instructions, if there are no cache misses, $S = 5$

Performance evaluation

- Pipelining improves performance by overlapping execution of successive instructions
 - Increases instruction throughput
 - Individual instruction is still executed in same number of cycles.
 - For five-stage pipeline, each instruction is executed in five cycles
 - But a new instruction can ideally enter the pipeline every cycle.

Performance evaluation

- In absence of stalls, $S = 1$, and ideal throughput with pipelining is:

$$P_p = f$$

- A 5-stage pipeline can potentially increase throughput by factor of five.
 - In general, an n -stage pipeline has the potential to increase throughput n times.

Effects of stalls/penalties

- Operations with longest delay dictate cycle time, and hence clock rate f .
- For a processor that has on-chip caches:
 - Memory-access operations have small delay when desired instructions or data are found in cache
- Delay through ALU is likely to be critical parameter.
 - If this delay is 2 ns, then $f = 500$ MHz
 - The ideal pipelined instruction throughput is

$$P_p = 500 \text{ MIPS (million instructions per second)}$$

Effects of stalls/penalties

- Consider a processor with operand forwarding in hardware
 - No penalties due to data dependencies, except in case of Load instructions.
- Consider how often a Load instruction is immediately followed by another instruction that uses result of memory access.
 - A one-cycle stall is necessary in such cases.
 - Ideal pipelined execution has $S = 1$, stalls due to such Load instructions have effect of increasing S by an amount δ_{stall} .

Effects of stalls/penalties

- Ex., Assume:
 - Load instructions constitute 25% of dynamic instruction count
 - 40% of Load instructions are followed by dependent instruction.
 - A one-cycle stall is needed in such cases.
 - Increase over the ideal case of $S = 1$ is
$$\delta_{\text{stall}} = 0.25 \times 0.40 \times 1 = 0.10$$
 - Execution time T is increased by 10 percent
 - Throughput is reduced to $P_p = f / (1 + \delta_{\text{stall}}) = f / 1.1 = 0.91f$

Effects of stalls/penalties

- Consider penalties due to mispredicting branches:
 - When both branch decision and branch target address are determined in Decode stage, branch penalty is one cycle.
 - Assume branches constitute 20% of dynamic instruction count, and average prediction accuracy for branch instructions is 90%
 - That is, 10% percent of all branch instructions that are executed incur a one-cycle penalty due to misprediction.
 - Increase in average number of cycles per instruction due to branch penalties is

$$\delta_{branch_penalty} = 0.20 \times 0.10 \times 1 = 0.02$$

Number of pipeline stages

- An n -stage pipeline:
 - May increase instruction throughput by a factor of n
 - Should use a large number of stages??
- As number of pipeline stages increases:
 - More instructions being executed concurrently.
 - More potential dependencies between instructions that may lead to pipeline stalls.
 - Branch penalty may be larger than one cycle if a longer pipeline moves the branch decision to a later stage.
- Gain in throughput from increasing value of n begins to diminish
 - Cost of a deeper pipeline may not be justified.

Number of pipeline stages

- Another important factor is inherent delay in basic operations performed by processor.
 - Most important is ALU delay.
 - In many processors, cycle time of processor clock is chosen so that one ALU operation can be completed in one cycle.
- Other operations, including accesses to a cache memory:
 - Typically divided into steps that each take about same time as an ALU operation.

Number of pipeline stages

- Further reductions in clock cycle time possible if a pipelined ALU is used.
 - Some recent processor implementations have used twenty or more pipeline stages to aggressively reduce the cycle time.

Superscalar Operations

- Equip processor with multiple execution units
 - Each may be pipelined, to increase processor's ability to handle several instructions in parallel.
 - Several instructions start execution in same clock cycle, but in different execution units
 - The processor is said to use multiple-issue.
 - Processors can achieve instruction execution throughput of more than one instruction per cycle.
 - Known as **superscalar processors**.
- Many modern high-performance processors use this approach.

Superscalar Operations

- To enable multiple-issue execution:
 - Superscalar processor has elaborate fetch unit
 - Fetches two or more instructions per cycle before they are needed
 - Places them in an instruction queue.
- A separate unit, called the dispatch unit:
 - Takes two or more instructions from front of queue
 - Decodes them, and sends them to appropriate execution units.
- At the end of the pipeline, another unit is responsible for writing results into register file.

Superscalar Operations

- A superscalar processor may incorporate two execution units
 - One for arithmetic instructions and another for Load and Store instructions.
 - Arithmetic operations normally require only one cycle, hence first execution unit is simple.
 - Load and Store instructions involve an address calculation for Index mode before each memory access, they have a two-stage pipeline

Superscalar Operations

- Challenges presented by branch instructions and data dependencies can be addressed with additional hardware.
 - Fetch unit handles branch instructions as it determines which instructions to place in queue for dispatching
 - It must determine both branch decision and target for each branch instruction
 - Branch decision may depend on result of an earlier instruction that is either still queued or newly dispatched

Superscalar Operations

- Stalling fetch unit until result is available can significantly reduce throughput
 - Not a desirable approach
 - Better to employ branch prediction

Superscalar Operations

- Since aim is to achieve high throughput
 - Prediction is combined with a technique called **speculative execution**.
 - Subsequent instructions based on an unconfirmed prediction are fetched, dispatched, and possibly executed
 - But labeled as being speculative so that they and their results may be discarded if the prediction is incorrect.

Superscalar Operations

- **Out-of-Order Execution:**

- Instructions are dispatched in same order as they appear in program.
- However, their execution may be completed out of order.

- Ex.,
 - Add R2, R3, #100
 - Load R5, 16(R6)
 - Subtract R7, R8, R9
 - Store R10, 24(R11)

- Subtract instruction writes to register R7 in same cycle as Load instruction that was fetched earlier writes to register R5.
- If memory access for Load instruction requires more than one cycle to complete, execution of Subtract instruction would be completed before Load instruction.

Superscalar Operations

- If an instruction I_{j+1} depends on the result of instruction I_j :
 - Execution of I_{j+1} will be delayed if result is not available when it is needed.
- As long as such dependencies are handled correctly:
 - No reason to delay execution of an unrelated instruction.
- If there is no dependency between a pair of instructions:
 - Order in which execution is completed does not matter
 - Results of execution of instructions must be written into destination locations strictly in program order

Superscalar Operations

- To improve performance, an execution unit allowed to execute any instructions whose operands are ready
 - May lead to out-of-order execution of instructions.
 - However, instructions must be completed in program order to allow precise exceptions.
- Can be resolved if execution is allowed to proceed out of order, but results are written into temporary registers.
 - Contents of these registers are later transferred to permanent registers in correct program order.
 - Last step is called **commitment** step, because effect of an instruction cannot be reversed after that point.

Superscalar Operations

- If an instruction causes an exception:
 - Results of any subsequent instructions that have been executed would still be in temporary registers and can be safely discarded.
 - Results that would normally be written to memory would also be buffered temporarily, and can be safely discarded as well.
- A temporary register that is assigned for result of an instruction assumes role of permanent register whose data it is holding.
 - Its contents are forwarded to any subsequent instruction that refers to original permanent register during that period.
 - This technique is called **register renaming**

Superscalar Operations

- When out-of-order execution is allowed:
 - Special control unit, called **commitment unit** is needed to guarantee in-order commitment.
 - Uses a separate queue called **reorder buffer** to determine which instruction(s) should be committed next.
 - Instructions are entered in queue strictly in program order as they are dispatched for execution.
- When instruction reaches head of this queue and execution of that instruction has been completed:
 - Corresponding results are transferred from temporary registers to permanent registers
 - Instruction is removed from queue.
 - All resources assigned to instruction, including temporary registers, are released.
 - Instruction is said to have been **retired** at this point.

Superscalar Operations

- Because an instruction is retired only when it is at the head of the queue:
 - All instructions that were dispatched before it must also have been retired.
 - Hence, instructions may complete execution out of order, but they are retired in program order.

END